

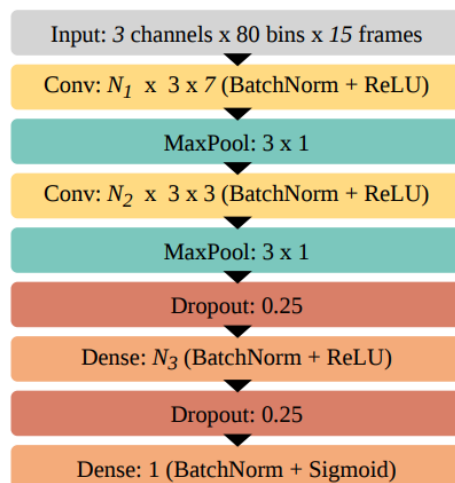
Stroke detection

Preprocessing:

- Audio Feature Extraction:
 - Spectrograms: Transform audio signals into spectrograms, providing a visual representation of frequency content over time.
 - Mel-Frequency Cepstral Coefficients (MFCCs): Extract coefficients representing the short-term power spectrum of sound, capturing crucial audio features.
 - Chroma Features: Describe pitch class energy distribution, useful for capturing tonal content.
- Data Augmentation:
 - Apply techniques like random pitch shifting, time stretching, and background noise addition to the training set for enhanced model robustness.

Model Architecture:

- Convolutional Neural Network (CNN) for Tabla Strokes:
 - Input Layer: Accept preprocessed audio features (e.g., spectrograms, MFCCs, chroma features).
 - Convolutional Layers: Extract hierarchical features, capturing patterns and relationships in audio features.
 - Pooling Layers: Downsample spatial dimensions, reducing computational load and enhancing translational invariance.
 - Flatten Layer: Flatten output from convolutional layers for input to fully connected layers.
 - Fully Connected Layers: Make predictions based on learned features, with the final layer having four nodes for tabla stroke categories.
 - Activation Functions: Use ReLU or similar functions to introduce non-linearity.



Transfer Learning:

- Source Domain (Western Drums):
 - Pretraining:
 - Utilize a dataset containing Western drum sounds.
 - Train initial CNN layers to capture general features related to drum sounds, initializing the network with cross-domain knowledge.
- Target Domain (Tabla Strokes):
 - Fine-Tuning:
 - Use the pretrained model as a starting point.
 - Replace or fine-tune later layers to adapt the model to tabla strokes.
 - Train the model on the tabla dataset to learn specific tabla stroke features.
 - Leverage knowledge gained from the source domain while adapting to tabla stroke nuances.

This holistic approach, combining detailed audio feature extraction, a robust CNN architecture, and strategic transfer learning, addresses challenges related to limited tabla stroke data and exploits similarities with Western drum sounds for improved model performance.

Basic code

```
import librosa
```

```
import librosa.display
```

```
import matplotlib.pyplot as plt
```

```
import numpy as np
```

```
import os
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import LabelEncoder
```

```
from tensorflow.keras.utils import to_categorical

# Function to load data

def load_data(data_path, labels_path)

    features = []

    labels = []

    for audio_file in os.listdir(data_path):

        if audio_file.endswith(".wav"):

            audio_path = os.path.join(data_path, audio_file)

            # Extract label from filename or use a separate file for labels

            label = get_label_from_filename(audio_file)

            labels.append(label)

        # Load audio file and extract features

        audio_data, _ = librosa.load(audio_path, sr=None)

        extracted_features = extract_features(audio_data)

        features.append(extracted_features)
```

```
return np.array(features), np.array(labels)
```

```
# Function to extract features from audio data
```

```
def extract_features(audio_data):
```

```
    # Using Mel-Frequency Cepstral Coefficients (MFCCs) as features
```

```
    mfccs = librosa.feature.mfcc(y=audio_data, sr=44100, n_mfcc=13)
```

```
    return mfccs.flatten()
```

```
# Function to apply data augmentation
```

```
def augment_data(train_features, train_labels):
```

```
    augmented_features = []
```

```
    augmented_labels = []
```

```
    for i in range(len(train_features)):
```

```
        augmented_audio = apply_augmentation(train_features[i])
```

```
        # Extract features from augmented audio
```

```
        augmented_features.append(extract_features(augmented_audio))
```

```
        augmented_labels.append(train_labels[i])
```

```

return np.array(augmented_features), np.array(augmented_labels)

# Function to apply random pitch shifting for data augmentation

def apply_augmentation(audio_data, sr=44100, pitch_shift_steps=2):

    # Randomly shift pitch within the specified number of steps

    pitch_shift_amount = np.random.uniform(-pitch_shift_steps, pitch_shift_steps)

    # Apply pitch shifting

    augmented_audio = librosa.effects.pitch_shift(audio_data, sr, n_steps=pitch_shift_amount)

    return augmented_audio

# Function to get label from filename (assuming filename is in the format "label_filename.wav")

def get_label_from_filename(filename):

    return filename.split("_")[0]

data_path = "path/to/audio/files"

labels_path = "path/to/labels"

features, labels = load_data(data_path, labels_path)

# Converting labels to one-hot encoding

```

```
label_encoder = LabelEncoder()
```

```
encoded_labels = label_encoder.fit_transform(labels)
```

```
one_hot_labels = to_categorical(encoded_labels)
```

```
# Split the dataset into training and testing sets
```

```
train_features, test_features, train_labels, test_labels = train_test_split(features, one_hot_labels,  
test_size=0.2, random_state=42)
```

Model

```
import tensorflow as tf
```

```
from tensorflow.keras import layers, models
```

```
def create_tabla_model(input_shape):
```

```
    model = models.Sequential()
```

```
    # Input Layer
```

```
    model.add(layers.InputLayer(input_shape=input_shape))
```

```
    # Convolutional Layers
```

```
    model.add(layers.Conv2D(32, (3, 3), activation='relu'))
```

```
    model.add(layers.MaxPooling2D((2, 2)))
```

```
    model.add(layers.Conv2D(64, (3, 3), activation='relu'))
```

```
model.add(layers.MaxPooling2D((2, 2)))

model.add(layers.Conv2D(128, (3, 3), activation='relu'))

model.add(layers.MaxPooling2D((2, 2)))

#Flatten Layer

model.add(layers.Flatten())

# Fully Connected Layers

model.add(layers.Dense(128, activation='relu'))

model.add(layers.Dropout(0.5))

model.add(layers.Dense(4, activation='softmax')) # 4 output nodes for the four tabla stroke
categories

# Compile the model

model.compile(optimizer='adam',

              loss='sparse_categorical_crossentropy',

              metrics=['accuracy'])

return model

# Assuming input_shape is the shape of your preprocessed audio features

input_shape = (num_frames, num_features, 1) # Update with your actual dimensions
```

```
tabla_model = create_tabla_model(input_shape)
```

```
# Display the model summary
```

```
tabla_model.summary()
```